

model_exploration

November 11, 2020

```
[1]: import os
import yaml

import librosa.display
import matplotlib.pyplot as plt
import numpy as np
import tensorflow as tf
```

```
[2]: ROOT = '/home/thomas/Dir/ccny/ccny-masters-thesis'
```

```
[3]: os.chdir(os.path.join(ROOT, 'tensorflow'))
from dataset import GE2EDatasetLoader
from loss import *
```

1 Load

In the process of training, we output two separate model objects. The first, is the **full** version, which simply uses the `model.save()` method as described [here](#). The second is the **embedding** model which has the cosine similarity layer at the end removed, thus its output is the d -vector. In the prediction step, all we care about is the eventual output of this embedding layer, which should serve as the speaker's *voiceprint* for a given utterance. When enrolling new speakers, typically we take a few utterances and *average* them to get their global voiceprint.

This is, however, not ultimately the model deployed on the edge device, as we must first compress it using **TFLite** (discussed below).

Here, we'll load the **embedding** model to get a sense of its performance on the task of speaker recognition and, more importantly, separation within the embedding space. We can also explore its capabilities with regards to *enrolling* new and unseen speakers.

```
[4]: embedding_models_path = os.path.join(ROOT, 'tensorflow/frozen_models/embedding')
full_models_path = os.path.join(ROOT, 'tensorflow/frozen_models/full')
```

```
[5]: os.listdir(embedding_models_path)
```

```
[5]: ['1605081214']
```

```
[6]: most_recent_embedding_model = max([ directory for directory in os.
    ↳listdir(embedding_models_path) ])
most_recent_full_model= max([ directory for directory in os.
    ↳listdir(full_models_path) ])
```

```
[7]: # print out the conf for this epoch
model_conf = os.path.join(ROOT, 'tensorflow/frozen_models/confs',
    ↳f'{most_recent_embedding_model}.yaml')
with open(model_conf, 'r') as stream:
    d = yaml.safe_load(stream)
d
```

```
[7]: {'sr': 16000,
'raw_data': {'path': '/media/thomas/TPD EX/thesis-data',
'datasets': {'train': {'LibriSpeech': ['train-clean-100',
'train-clean-360',
'train-other-500'],
'VoxCeleb1': ['dev']}},
'test': {'out_of_sample': {'LibriSpeech': ['test-clean', 'test-other']}}},
'feature_data': {'path': '/home/thomas/Dir/ccny/ccny-masters-thesis/feature-
data',
'speakers_per_batch': 8,
'utterances_per_speaker': 8},
'features': {'type': 'melspectrogram',
'window_length': 1.2,
'overlap_percent': 0.5,
'frame_length': 0.025,
'hop_length': 0.01,
'n_fft': 512,
'n_mels': 40,
'trim_top_db': 20},
'train': {'epochs': 50,
'network': {'optimizer': {'type': 'SGD', 'lr': 0.01, 'clipnorm': 3.0},
'dropout': 0.1,
'layers': [{'lstm': {'units': 128, 'return_sequences': True}},
{'lstm': {'units': 128}},
{'embedding': {'nodes': 128}},
{'similarity_matrix': {'embedding_length': 128}}],
'callbacks': {'lr_scheduler': {'cutoff_epoch': 25, 'decay': 'exponential'},
'csv_logger': {'dir': 'training_logs'},
'checkpoint': {'dir': 'model_checkpoints'}}}}
```

```
[8]: embedding_model = tf.keras.models.load_model(os.path.
    ↳join(embedding_models_path, most_recent_embedding_model))
embedding_model.summary()
```

WARNING:tensorflow:No training configuration found in save file, so the model

was **not** compiled. Compile it manually.

[WARNING] {tensorflow} 2020-11-11 19:23:32,613 No training configuration found in save file, so the model was **not** compiled. Compile it manually.

Model: "sequential_1"

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, 40, 128)	128000
lstm_1 (LSTM)	(None, 128)	131584
dense (Dense)	(None, 128)	16512
Total params: 276,096		
Trainable params: 276,096		
Non-trainable params: 0		

```
[9]: full_model = tf.keras.models.load_model(os.path.join(full_models_path,
    ↪most_recent_full_model), compile=False)
full_model.compile(
    loss=get_embedding_loss(N=8, M=8)
)
full_model.summary()
```

Model: "speaker_verification_model"

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, 40, 128)	128000
lstm_1 (LSTM)	(None, 128)	131584
dense (Dense)	(None, 128)	16512
speaker_similarity_matrix_la (64, 8)		2
sequential (Sequential)	(64, 8)	276098
Total params: 276,098		
Trainable params: 276,098		
Non-trainable params: 0		

2 Example Predict


Now that we have de-serialized our model into `m`, it's only useful if we can make predictions. Here we'll demonstrate how to use the `__call__` method to generate embedding for a batch from the training dataset. Notice, the first dimension of each layer is `None`, which corresponds to our `batch_size`. This means we should be able to pass a feature of arbitrary batches, e.g. $[N, h, w]$ where N is the number of samples, h is the height of the spectrogram feature and w is the width of the spectrogram feature. This is preferable than needing to batch examples in a specific number at inference time.

```
[10]: dataset = GE2EDatasetLoader(  
        root_dir=os.path.join(ROOT, 'feature-data')  
    )  
    metadata = dataset.get_metadata()  
    inputs, targets = dataset.get_single_train_batch()
```

```
[11]: inputs.shape
```

```
[11]: TensorShape([64, 40, 121])
```

```
[12]: # here's a sample feature: Mel-based spectrogram  
    plt.figure(figsize=(8,4))  
    _ = librosa.display.specshow(inputs[0].numpy())
```



```
[13]: embedding_length = embedding_model.layers[-1].output.shape[1]  
    embedding_length
```

```
[13]: 128
```

```
[14]: embeddings = embedding_model(inputs)
      embeddings.shape # [batch_size x embedding_length]
```

```
[14]: TensorShape([64, 128])
```

```
[15]: # how to visualize?
```

3 Speaker separation

Now that we've demonstrated how to generate predictions, we'll explore how our model performs at its primary task: separating unique speakers within the embedding space. Given that we're working in a high dimensional space, we'll need to use techniques that will allow us humans to actually *understand* that separation.

```
[16]: train_dataset, test_dataset = dataset.get_datasets()
      train_it, test_it = iter(train_dataset), iter(test_dataset)

      metadata = dataset.get_metadata()
      speaker_id_mapping = { v:k for k, v in metadata['speaker_id_mapping'].items() }
```

```
[17]: def get_orig_speakers(tensor, mapping):
      return [mapping[t] for t in tensor.numpy()]
```

3.0.1 t-SNE

In order to *visualize* high dimensional data we need to make it intelligible for humans. Here we'll attempt to gain an understanding of the separation in our embedding space by using the [t-SNE](#) procedure. Here, points in our high dimensional space which are *similar* (i.e. their dot product is high) should be grouped together.

```
[18]: from sklearn.manifold import TSNE
```

```
[19]: batch_size = metadata['batch_size']
      n_samples = batch_size * 1 # batch_size x n_batches

      inputs = np.zeros((n_samples, embedding_length))
      targets = np.zeros((n_samples,))


      i = 0
      while i < n_samples:
          x, y = next(train_it)
          embeddings = embedding_model(x)

          inputs[i:i+batch_size,:] = embeddings
          targets[i:i+batch_size] = y
```

```
i += batch_size
```

```
[79]: X_embedded = TSNE(n_components=2).fit_transform(inputs)

plt.figure(figsize=(10,6))
_ = plt.scatter(X_embedded[:, 0], X_embedded[:, 1], c=targets, cmap='seismic')
```



3.1 Cosine similarity


Similarly, we can visualize the similarity of embedding vectors for each utterance in a batch. This should yield a block diagonal matrix where each $N \times M$ block along the diagonal (the i^{th} speaker's embeddings) has a higher cosine similarity than the rest of the embedding vectors.

```
[80]: from sklearn.metrics.pairwise import cosine_similarity
```

```
[81]: cos = cosine_similarity(embeddings)
cos.shape
```

```
[81]: (64, 64)
```

```
[82]: plt.figure(figsize=(12,8))
_ = plt.imshow(cos, cmap='hot', interpolation='nearest')
```



4 Validation

4.1 Out-of-Sample Dataset

```
[24]: full_model.evaluate(dataset.get_test_dataset())
```

288/288 [=====] - 54s 188ms/step - loss: 0.1659

```
[24]: 0.16587024927139282
```

4.2 Equal Error Rate (EER)

In many biometric security systems the [equal error rate](#) gives us a performance idea about the model. The equal error ratio depends on two statistics the false acceptance ratio, FAR , and the false rejection ratio, FRR .

FAR is the ratio of falsely accepted scored impostors over the total scored impostors, e.g. the number of utterances which should have been rejected as being close to a true speaker d -vector.

Conversely, the FRR are the utterances which should have been close to a true speaker d -vector, but weren't.

We find the EER when the average of these two is at its minimum for some threshold t .

```
[25]: eer_inputs, _ = dataset.get_single_test_batch()
     eer_preds = full_model(eer_inputs)
     eer_pred_norm = tf.math.l2_normalize(eer_preds, axis=0)
```

```
[28]: threshold, eer, far, frr = equal_error_ratio(eer_pred_norm, 8, 8, 0.0)
     threshold, eer
```

```
[28]: (0.26000000000000006, 0.1328125)
```

```
[41]: # calculate averages
count, sum_thres, sum_eer = 0, 0.0, 0.0
for x, _ in dataset.get_test_dataset():
    S = full_model(eer_inputs)
    S_norm = tf.math.l2_normalize(S, axis=0)
    threshold, eer, _, _ = equal_error_ratio(S_norm, 8, 8, 0.0)
    sum_thres += threshold
    sum_eer += eer
    count += 1
```

```
[42]: # avg threshold
     sum_thres / count
```

```
[42]: 0.26000000000000002
```

```
[43]: # avg EER
     sum_eer / count
```

```
[43]: 0.1328125
```

5 Enrollment

As previously stated, the generalized approach is intended to demonstrate that our model can *learn* how to separate speakers in a high dimensional embedding space such that we can distinguish utterances spoken by different people. This is useful in a whole bunch of tasks that require biometric

security, personalization, etc. Just think of how your Google Home or Amazon Alexa is able to tell *which* person in your household is talking to it. This is how.

So, we'll walk through *enrolling* a new speaker. Here that means passing some of their samples through our model and then averaging out the embedding vector to get their *d*-vector. This will then serve as the reference point for that person. When we get a **new** utterance, we will check it against that *d*-vector in order to accept/reject based on some threshold.

```
[44]: enrollment_inputs, enrollment_targets = dataset.get_single_test_batch()
      enrollment_speakers = np.unique(enrollment_targets)
```

```
[45]: idx = 0
      speaker = enrollment_speakers[idx]
      speaker
```

```
[45]: 3558
```

```
[46]: orig_id = None
      for full_id, mapped_id in dataset.get_metadata()['speaker_id_mapping'].items():
          if mapped_id == speaker:
              orig_id = full_id
      orig_id
```

```
[46]: 'LibriSpeech/test-clean/8463'
```

```
[47]: # how many utterances do we have for this person?
      indices = [ i for i, s in enumerate(enrollment_targets) if s == speaker]

      begin, end = indices[0], indices[len(indices)//2]
      len(indices), indices
```

```
[47]: (8, [40, 41, 42, 43, 44, 45, 46, 47])
```

```
[48]: # now create an embedding vector for this person from end-begin=4 utterances_
      ↪above
      enrollment_embeddings = embedding_model(enrollment_inputs)

      d_vector = np.mean(enrollment_embeddings[begin:end], axis=0)
      d_vector.shape
```

```
[48]: (128,)
```

The metric with which we compare each enrolled *d*-vector and a new *d*-vector for an unseen utterance is the cosine similarity, which we explained above.

```
[49]: # find cosine similarity with same speaker
      cosine_similarity(d_vector.reshape(1, -1), enrollment_embeddings[end+2].numpy().
      ↪reshape(1, -1))
```

```
[49]: array([[0.74438655]], dtype=float32)
```

```
[50]: # with utterance of a different speaker  
cosine_similarity(d_vector.reshape(1, -1), enrollment_embeddings[0].numpy().  
↳reshape(1, -1))
```

```
[50]: array([[0.14067599]], dtype=float32)
```