

tfLite_conversion

November 18, 2020

```
[5]: import os
import sys
import yaml

import numpy as np
from sklearn.metrics.pairwise import cosine_similarity
import tensorflow as tf
```

```
[6]: tf.__version__
```

```
[6]: '2.3.1'
```

```
[7]: ROOT = '/home/thomas/Dir/ccny/ccny-masters-thesis'
```

```
[8]: sys.path.append(os.path.join(ROOT, 'tensorflow'))

from dataset import GE2EDatasetLoader
```

1 Load tf.keras Model

```
[9]: embedding_models_path = os.path.join(ROOT, 'tensorflow/frozen_models/embedding')
model_confs_path = os.path.join(ROOT, 'tensorflow/frozen_models/confs')
```

```
[10]: models = os.listdir(embedding_models_path)
models
```

```
[10]: ['1605081214',
      '1605371725',
      '1605623350',
      '1605672336',
      '1605647312',
      '1605518238']
```

```
[11]: model_idx = 3

model_to_convert = models[model_idx]
```

```
model_to_convert
```

```
[11]: '1605672336'
```

```
[12]: confs = os.listdir(os.path.join(model_confs_path, model_to_convert))
      confs
```

```
[12]: ['conv_1d.yml']
```

```
[13]: with open(f'{model_confs_path}/{model_to_convert}/{confs[0]}', 'r') as c:
      conf = yaml.safe_load(c)

      conf['train']['network']
```

```
[13]: {'optimizer': {'type': 'SGD', 'lr': 0.01, 'clipnorm': 3.0},
      'dropout': 0.1,
      'layers': [{'conv1d': {'filters': 8, 'kernel_size': 10}},
                  {'conv1d': {'filters': 8, 'kernel_size': 3}},
                  'flatten',
                  {'embedding': {'nodes': 64}},
                  {'similarity_matrix': {'embedding_length': 64}}],
      'callbacks': {'lr_scheduler': {'cutoff_epoch': 25, 'decay': 'exponential'},
                    'csv_logger': {'dir': 'training_logs'},
                    'checkpoint': {'dir': 'model_checkpoints'}}
```

```
[14]: model_path = os.path.join(embedding_models_path, model_to_convert)
```

```
[15]: embedding_model = tf.keras.models.load_model(model_path)
      embedding_model.summary()
```

WARNING:tensorflow:No training configuration found in save file, so the model was *not* compiled. Compile it manually.

[WARNING] {tensorflow} 2020-11-18 07:04:37,060 No training configuration found in save file, so the model was *not* compiled. Compile it manually.

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv1d (Conv1D)	(None, 31, 8)	9688
batch_normalization (Batch Normalization)	(None, 31, 8)	32
dropout (Dropout)	(None, 31, 8)	0
conv1d_1 (Conv1D)	(None, 29, 8)	200

batch_normalization_1 (Batch Normalization)	(None, 29, 8)	32

dropout_1 (Dropout)	(None, 29, 8)	0

flatten (Flatten)	(None, 232)	0

dense (Dense)	(None, 64)	14912
=====		
Total params: 24,864		
Trainable params: 24,832		
Non-trainable params: 32		

```
[16]: train, test = GE2EDatasetLoader(root_dir=os.path.join(ROOT, 'feature-data')).
      ↪get_datasets()
```

2 Convert to tf.lite model

```
[42]: converter = tf.lite.TFLiteConverter.from_saved_model(model_path)
      converter.optimizations = [tf.lite.Optimize.DEFAULT]
      tflite_model = converter.convert()
```

```
[43]: TFLITE_MODEL_DIR = os.path.join(ROOT, f'tensorflow/frozen_models/tiny/
      ↪{model_to_convert}')
      TFLITE_MODEL_PATH = os.path.join(ROOT, TFLITE_MODEL_DIR, 'model.tflite')
      !mkdir -p $TFLITE_MODEL_DIR
```

```
[44]: with open(TFLITE_MODEL_PATH, 'wb') as out:
      out.write(tflite_model)
```

```
[45]: !du -sh $TFLITE_MODEL_DIR/*
```

```
32K      /home/thomas/Dir/ccny/ccny-masters-
thesis/tensorflow/frozen_models/tiny/1605672336/model.tflite
```

3 Making predictions

```
[46]: feats, targets = next(iter(train))
```

```
[47]: feats.shape # [batch_size x spectrogram height x spectrogram width]
```

```
[47]: TensorShape([64, 40, 121])
```

```
[48]: targets[:8].numpy() # 8 utterances / per speaker / per batch
```

```
[48]: array([2067, 2067, 2067, 2067, 2067, 2067, 2067, 2067])
```

```
[49]: def get_output(feat, input_details, output_details):  
    if feat.ndim == 2:  
        feat = np.expand_dims(feat, axis = 0)  
        feat = feat.astype(input_details["dtype"])  
        interpreter.set_tensor(input_details["index"], feat)  
        interpreter.invoke()  
        output = interpreter.get_tensor(output_details["index"])[0]  
    return output
```

```
[50]: interpreter = tf.lite.Interpreter(model_path=str(TFLITE_MODEL_PATH))  
interpreter.allocate_tensors()
```

```
[51]: input_details = interpreter.get_input_details()[0]  
input_details
```

```
[51]: {'name': 'conv1d_input',  
      'index': 0,  
      'shape': array([ 1, 40, 121], dtype=int32),  
      'shape_signature': array([-1, 40, 121], dtype=int32),  
      'dtype': numpy.float32,  
      'quantization': (0.0, 0),  
      'quantization_parameters': {'scales': array([], dtype=float32),  
      'zero_points': array([], dtype=int32),  
      'quantized_dimension': 0},  
      'sparsity_parameters': {}}
```

```
[52]: output_details = interpreter.get_output_details()[0]  
output_details
```

```
[52]: {'name': 'Identity',  
      'index': 27,  
      'shape': array([ 1, 64], dtype=int32),  
      'shape_signature': array([-1, 64], dtype=int32),  
      'dtype': numpy.float32,  
      'quantization': (0.0, 0),  
      'quantization_parameters': {'scales': array([], dtype=float32),  
      'zero_points': array([], dtype=int32),  
      'quantized_dimension': 0},  
      'sparsity_parameters': {}}
```

```
[53]: # d-vector  
  
# enroll with 4 utterances  
n_enrollment = 4
```

```

d_vector = np.zeros((output_details['shape'][1],))

print(f'enrolling speaker_id: {np.unique(targets[:n_enrollment])}')
for i in range(n_enrollment):
    d_vector += get_output(feats[i], input_details, output_details)

d_vector /= n_enrollment
d_vector.shape

```

enrolling speaker_id: [2067]

[53]: (64,)

```

[56]: # test against another example from same speaker
speaker_idx = n_enrollment + 1

pred = get_output(feats[speaker_idx], input_details, output_details)
print(f'speaker_id: {targets[speaker_idx]}')
cosine_similarity(d_vector.reshape(1,-1), pred.reshape(1,-1))

```

speaker_id: 2067

[56]: array([[0.81528801]])

```

[57]: # test against utterance from different speaker
speaker_idx = n_enrollment + 5

pred = get_output(feats[speaker_idx], input_details, output_details)
print(f'speaker_id: {targets[speaker_idx]}')
cosine_similarity(d_vector.reshape(1,-1), pred.reshape(1,-1))

```

speaker_id: 2867

[57]: array([[0.30849305]])

4 Convert to C++ array

The real kicker here is being able to port our `tf.lite` model onto our target microcontroller (in this case an [Arduino Nano 33 BLE](#)). With the model represented on our microcontroller, we can do online speaker verification at the edge in a low power + compute setting.

```

[58]: # https://github.com/tensorflow/tensorflow/tree/master/tensorflow/lite/micro/
      ↪ examples/micro\_speech

```

[59]:

```
TFLITE_MODEL_MICRO_PATH = os.path.join(ROOT, 'MicrocontrollerRecognizer/model.
↳h')
TFLITE_MODEL_MICRO_PATH
```

```
[59]: '/home/thomas/Dir/ccny/ccny-masters-thesis/MicrocontrollerRecognizer/model.h'
```

```
[60]: # https://colab.research.google.com/github/tensorflow/tensorflow/blob/master/
↳tensorflow/lite/micro/examples/hello_world/train/train_hello_world_model.
↳ipynb#scrollTo=_UQbInrLd_ET
```

```
[61]: # Update variable names
!xxd -i {TFLITE_MODEL_PATH} > {TFLITE_MODEL_MICRO_PATH}
REPLACE_TEXT = TFLITE_MODEL_PATH.replace('/', '_').replace('.', '_').
↳replace('-', '_')
!sed -i 's/{REPLACE_TEXT}/speaker_model/g' {TFLITE_MODEL_MICRO_PATH}
```

```
[62]: !head {TFLITE_MODEL_MICRO_PATH}
```

```
unsigned char speaker_model[] = {
    0x28, 0x00, 0x00, 0x00, 0x54, 0x46, 0x4c, 0x33, 0x00, 0x00, 0x00, 0x00,
    0x00, 0x00, 0x00, 0x00, 0x00, 0x00, 0x00, 0x00, 0x00, 0x00, 0x12, 0x00,
    0x1c, 0x00, 0x04, 0x00, 0x08, 0x00, 0x0c, 0x00, 0x10, 0x00, 0x14, 0x00,
    0x00, 0x00, 0x18, 0x00, 0x12, 0x00, 0x00, 0x00, 0x03, 0x00, 0x00, 0x00,
    0x3c, 0x7b, 0x00, 0x00, 0x80, 0x67, 0x00, 0x00, 0x68, 0x67, 0x00, 0x00,
    0x3c, 0x00, 0x00, 0x00, 0x04, 0x00, 0x00, 0x00, 0x01, 0x00, 0x00, 0x00,
    0x0c, 0x00, 0x00, 0x00, 0x08, 0x00, 0x0c, 0x00, 0x04, 0x00, 0x08, 0x00,
    0x08, 0x00, 0x00, 0x00, 0x08, 0x00, 0x00, 0x00, 0x1d, 0x00, 0x00, 0x00,
    0x13, 0x00, 0x00, 0x00, 0x6d, 0x69, 0x6e, 0x5f, 0x72, 0x75, 0x6e, 0x74,
```

```
[ ]:
```